

# ECE60146 DL HW2

Akshita Kamsali, akamsali@purdue.edu

January 24, 2023

Note: bold lower case letters indicate vectors and bold upper case letters indicate matrices

## 1 Theory

The mystery of how converting to float and dividing by max value yields the same result as `ToTensor` from `torchvision.transforms` can be solved from looking at the source code for the implementation of latter.

In fig 1 we can see that the image is reshape to  $\mathbf{C} \times \mathbf{H} \times \mathbf{W}$  and then divided by 255 to normalise. This is based on the assumption that PIL images are 8 byte integers of shape  $\mathbf{H} \times \mathbf{W}$  with  $\mathbf{C}$  channel number of tuples at every position. Since the maximum value an 8 byte integer can take is 255, it is hardcoded as such.

## 2 Programming

### 2.1 Task 1: Applying Affine/Perspective transform

In the field computer vision, homography is used a fundamental tool to perform many image based tasks. When training a neural network, we want the model to learn to identify these projective distortions and illumination differences. To show one such transformation we pick points from a front view image of a stop sign and oblique view image. Then we try to transform the oblique view into a front view with either `RandomAffine` or `perspective` functions available to us in `torchvision.transforms`.

Figure 2 shows `RandomAffine` and perspective transform to tranform oblique view to front view. Figures 3 and 4 show zoomed version from 2b and 2c respectively.

```
151
152     img = torch.from_numpy(pic.transpose((2, 0, 1))).contiguous()
153     # backward compatibility
154     if isinstance(img, torch.ByteTensor):
155         return img.to(dtype=default_float_dtype).div(255)
156     else:
157         return img
158
```

Figure 1: Source code from PyTorch GitHub Repo

Method	Runtime
MyDataset	635.8
Dataloader ( $batchsize = 4, nw = 1$ )	672
Dataloader ( $batchsize = 4, nw = 2$ )	415
Dataloader ( $batchsize = 4, nw = 4$ )	311
Dataloader ( $batchsize = 8, nw = 2$ )	350
Dataloader ( $batchsize = 8, nw = 4$ )	300

Table 1: Runtimes for my implementation and Pytorch Dataloader with varying workers (nw)

**Similarity measurement** Then we generate histograms from two images to measure similarity using wasserstain distance. Wasserstein distance between the two images is 0.001020476357306882. Histograms are shown in 5

## 2.2 Task 2: Using torch Dataloader

**Resize and Normalise** I chose to resize images to (1024, 1024), followed by conversion to tensors using `ToTensor()`, then `Normalise`.

**Gaussian blur to enhance features at scales** After that I made the design choice of using `Gaussianblur`. I made this choice to enhance image structures at different scales, and simultaneously reduce noise.

**Affine and Perspective transforms** Then I chose to apply `Affine` and `Perspective` transforms to introduce affine and projective tranformations at random to the images in dataset. This to done to augment more types of scenarios where the images may be taken from oblique angles and distorted due to varying perceptions.

### 2.2.1 Comparision of `__getitem__` and Parallel Dataloader

Time taken for 1000 calls of `__getitem__` 635.8037867546082.

now with batch size fixed at 4:

Time taken for 1000 calls of torch Dataloader with one worker: 671.9216966629028

Time taken for 1000 calls of torch Dataloader with two workers: 415.3375608921051

Time taken for 1000 calls of torch Dataloader with  $nw = 4$ : 311.74111437797546

Table 1 shows runtimes.

I observe that absolute computation times vary depending on hardware. However, the pattern looks very similar with computation time reducing with increasing batch size and number of workers. Figures 6 and 7 show images from `__getitem__` and torch Dataloader calls.



(a) Front view and Oblique view



(b) Front view and Oblique view with Affine transformed applied



(c) Front view and Oblique view with Affine and perspective transformed applied

Figure 2: Task 1 figures



Figure 3: Oblique view with Affine transformation



Figure 4: Oblique view with Affine followed by perspective transformation

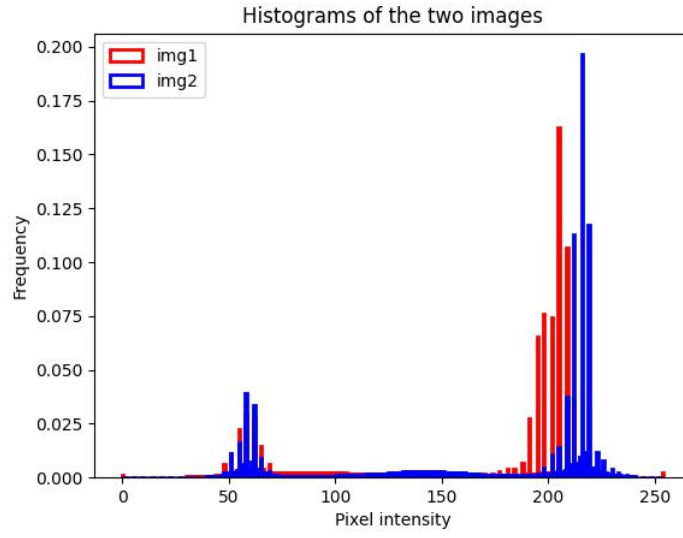


Figure 5: Histogram showing histograms for two different perspectives after transformation

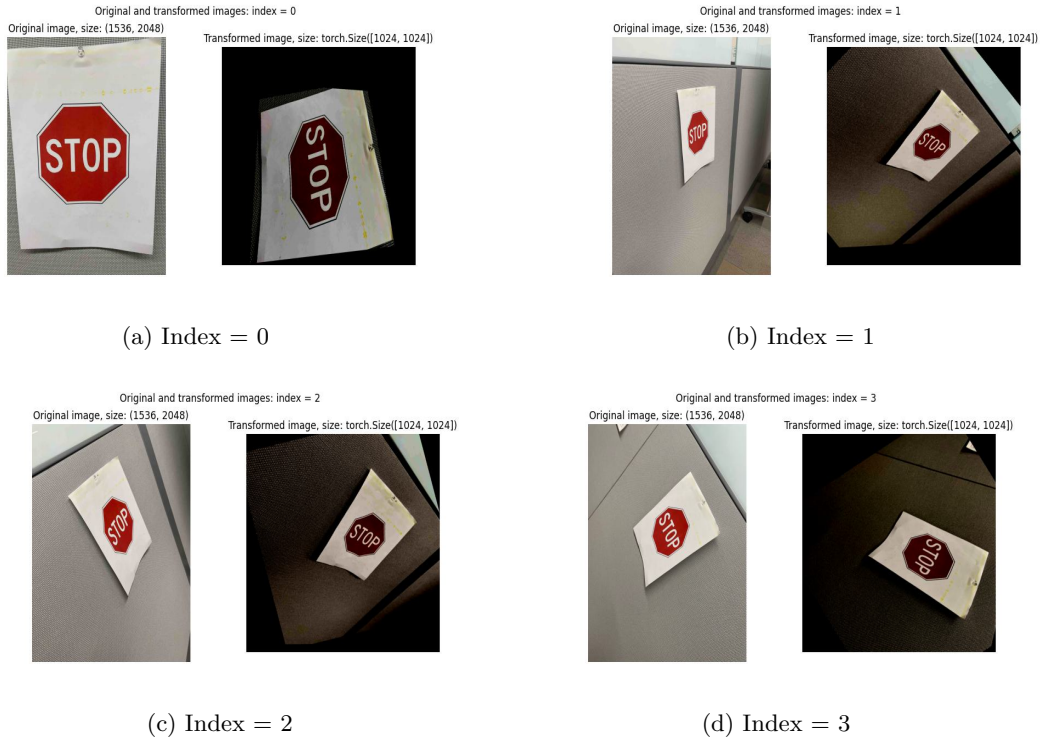
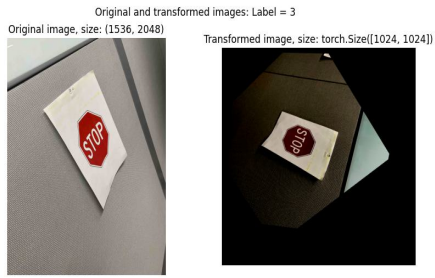
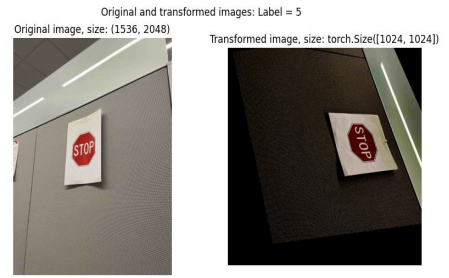


Figure 6: Plots with 4 calls to `__getitem__`, image before and after transformation



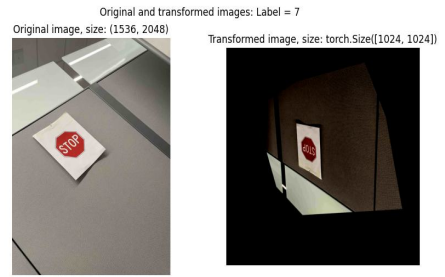
(a) Label = 3



(b) Label = 5



(c) Label = 10



(d) Label = 7

Figure 7: Plots with iterating over over first batch of the Dataloader with batchsize 4

## 3 Code

I have written my code for `MyDataset` in a separate .py file and call it from my main file.

### 3.1 MyDataset

```
1 import glob
2 from PIL import Image
3 from torchvision import transforms as tvt
4 import torch
5
6
7 class MyDataset(torch.utils.data.Dataset):
8     def __init__(self, root: "str") -> None:
9         super().__init__()
10        # get file list with glob.glob
11        self.file_list = sorted(glob.glob(root + "/*.jpg"))
12        # set transform with tvt.Compose
13        self.transform = tvt.Compose(
14            [
15                tvt.Resize((1024, 1024)),
16                tvt.ToTensor(),
17                tvt.Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5]),
18                tvt.GaussianBlur(3, sigma=(0.1, 2.0)),
19                tvt.RandomAffine(
20                    degrees=(0, 90),
21                    translate=(0.1, 0.1),
22                    scale=(0.8, 1.2),
23                    shear=10,
24                    fill=0,
25                ),
26                tvt.RandomPerspective(distortion_scale=0.5, p=0.5, fill=0),
27            ]
28        )
29
30    def __len__(self) -> int:
31        return len(self.file_list)
32
33    def __getitem__(self, index: int):
34        # read image with PIL.Image.open
35        img = Image.open(self.file_list[index])
36        # transform image with self.transform
37        img = self.transform(img)
38        label = self.file_list[index].split("/")[1].split(".")[0]
39
40        return img, int(label)
```

mydataset.py

### 3.2 main

```
1 from mydataset import MyDataset
2
3 import torch
4 from torch.utils.data import DataLoader
5 import torchvision.transforms as tvt
6
```

```

7 from scipy.stats import wasserstein_distance
8 from PIL import Image
9
10 import numpy as np
11 import matplotlib.pyplot as plt
12 from time import time
13
14 import matplotlib.pyplot as plt
15
16 img1 = Image.open("data/1.jpg")
17 img2 = Image.open("data/2.jpg")
18
19
20 def concatenate_img(img1, img2, name="concat.jpg"):
21     concat_img = Image.new("RGB", (img1.size[0] * 2, img1.size[1]), "white")
22
23     # pasting the first image (image_name, (position))
24     concat_img.paste(img1, (0, 0))
25     # pasting the second image (image_name, (position))
26     concat_img.paste(img2, (img1.size[0], 0))
27
28     # save the concatenated image
29     plt.imshow(concat_img)
30     plt.axis("off")
31     plt.savefig(f"solutions/{name}")
32
33
34 concatenate_img(img1, img2)
35
36
37 # apply affine transform to img2
38 affine_tr = tvt.RandomAffine((30, 40))
39 img2 = affine_tr(img2)
40
41 # save the transformed image
42 img2.save("solutions/2_aff.jpg")
43 concatenate_img(img1, img2, name="concat_aff.jpg")
44
45
46 # apply perspective transform to img2
47 start_pts = [[851, 728], [973, 683], [851, 1171], [961, 1110]]
48 end_pts = [[546, 572], [917, 578], [579, 1431], [925, 1431]]
49
50 transformed_image = tvt.functional.perspective(
51     img2, startpoints=start_pts, endpoints=end_pts
52 )
53 transformed_image.save("solutions/2_aff_persp.jpg")
54 concatenate_img(img1, transformed_image, name="concat_aff_persp.jpg")
55
56
57 # generate histograms
58 def get_hist(img, bins=10):
59     tensor_img = tvt.functional.pil_to_tensor(img.convert('L')).float()
60     hist_img = torch.histc(tensor_img, bins=bins, min=0, max=255)
61     hist_img = hist_img.div(hist_img.sum()) # normalize the histogram for probability
62     return hist_img
63
64 # get histograms of the two images
65 bins = 255
66 hist_img1 = get_hist(img1, bins=bins)

```



```

67 hist_img2 = get_hist(transformed_image, bins=bins)
68
69 # plot the histograms
70
71 plt.bar(np.arange(bins), hist_img1, fill=False, edgecolor='red', linewidth=2, label='
img1')
72 plt.bar(np.arange(bins), hist_img2, fill=False, edgecolor='blue', linewidth=2, label=
'img2')
73 plt.legend()
74 plt.title('Histograms of the two images')
75 plt.ylabel('Frequency')
76 plt.xlabel('Pixel intensity')
77 plt.savefig('solutions/hist.jpg')
78
79
80 # compute wasserstein distance
81 dist = wasserstein_distance(hist_img1.cpu().numpy(),
82                             hist_img2.cpu().numpy())
83 print(f'Wasserstein distance between the two images is {dist}')
84
85
86 my_dataset = MyDataset('data')
87 print(len(my_dataset))
88 index = 9
89 print(my_dataset[index][0].shape, my_dataset[index][1])
90
91 def plot Og_tr_img(og_img, tr_img, index=9):
92     fig, ax = plt.subplots(1, 2, figsize=(10, 5))
93     ax[0].imshow(og_img)
94     ax[0].axis('off')
95     ax[0].set_title(f'Original image, size: {og_img.size}')
96     ax[1].imshow(tr_img)
97     ax[1].set_title(f'Transformed image, size: {tr_img.size()[2]}')
98     ax[1].axis('off')
99     plt.suptitle(f'Original and transformed images: Label = {index}')
100    plt.savefig(f'solutions/dl Og_tr_img-{index}.jpg')
101
102 # plot images from getitem
103 for i in range(10):
104     og_img = Image.open(my_dataset.file_list[i])
105     tr_img = my_dataset[i][0].permute(1,2,0)
106     plot Og_tr_img(og_img, tr_img, index=i)
107
108 # plot images from dataloader
109 batch_size = 4
110 train_dataloader = DataLoader(my_dataset, batch_size=batch_size, shuffle=True,
    num_workers=2)
111
112 train_data, labels = next(iter(train_dataloader))
113
114 #plot images from a batch in train dataloader
115 for i in range(batch_size):
116     og_img = Image.open(my_dataset.file_list[labels[i].item() - 1])
117     tr_img = train_data[i].permute(1,2,0)
118     plot Og_tr_img(og_img, tr_img, index=labels[i])
119
120
121
122 # Time taken for 1000 calls of --getitem--
123 start = time()

```

```

124 num_tot_images = len(my_dataset)
125 for i in range(1000):
126     my_dataset.__getitem__(i%num_tot_images)
127
128 print("Time taken for 1000 calls of __getitem__", time() - start)
129
130 # Time taken for 1000 image accumulations in torch DataLoader
131 batch_size = 4
132 train_dataloader = DataLoader(
133     my_dataset, batch_size=batch_size, shuffle=True, num_workers=2
134 ) # change num_workers to 0/1/4 to see the difference
135
136 start = time()
137
138 count = 0
139 while True:
140     for img, labels in train_dataloader:
141         count += labels.shape[0]
142     if count >= 1000:
143         break
144
145 print("Time taken for 1000 calls of torch DataLoader", time() - start)

```

main.py